

CS 581 – ADVANCED ARTIFICIAL INTELLIGENCE

TOPIC: DECISION MAKING UNDER UNCERTAINTY



Mustafa Bilgic



<http://www.cs.iit.edu/~mbilgic>

MOTIVATION

- The agent needs reason in an uncertain world
- Uncertainty can be due to
 - Noisy sensors (e.g., temperature, GPS, camera, etc.)
 - Imperfect data (e.g., low resolution image)
 - Missing data (e.g., lab tests)
 - Imperfect knowledge (e.g., medical diagnosis)
 - Exceptions (e.g., all birds fly except ostriches, penguins, birds with injured wings, dead birds, ...)
 - Changing data (e.g., flu seasons, traffic conditions during rush hour, etc.)
 - ...
- The agent still must act (e.g., step on the breaks, diagnose a patient, order a lab test, ...)

WHAT TO DO?

- Probability of a given email being spam is 0.6. Where should that email be delivered?
- Probability of rain is 0.6. I don't like carrying my umbrella around, but I also don't like getting wet. Should I take my umbrella?
- Probability of a given patient suffering from a heart disease is 0.6. Should more tests be conducted and if so, which ones and in which order? What should the treatment plan be?
- Probability of the product selling is 0.6. If we spend \$100K in ads, the probability goes up to 0.7. Is \$100K on ads justified?

RATIONAL AGENT

- Given world states, a utility function, actions, transitions, evidence, and probabilities,
- A rational agent chooses the action that maximizes expected utility

$$action = \operatorname{argmax}_a EU(a|e)$$

UNCERTAINTY REPRESENTATION

○ Covered

- Probability background (marginal, conditional, independence, Bayes rule, ...)
- Probabilistic classification (naïve Bayes, logistic regression, neural networks)

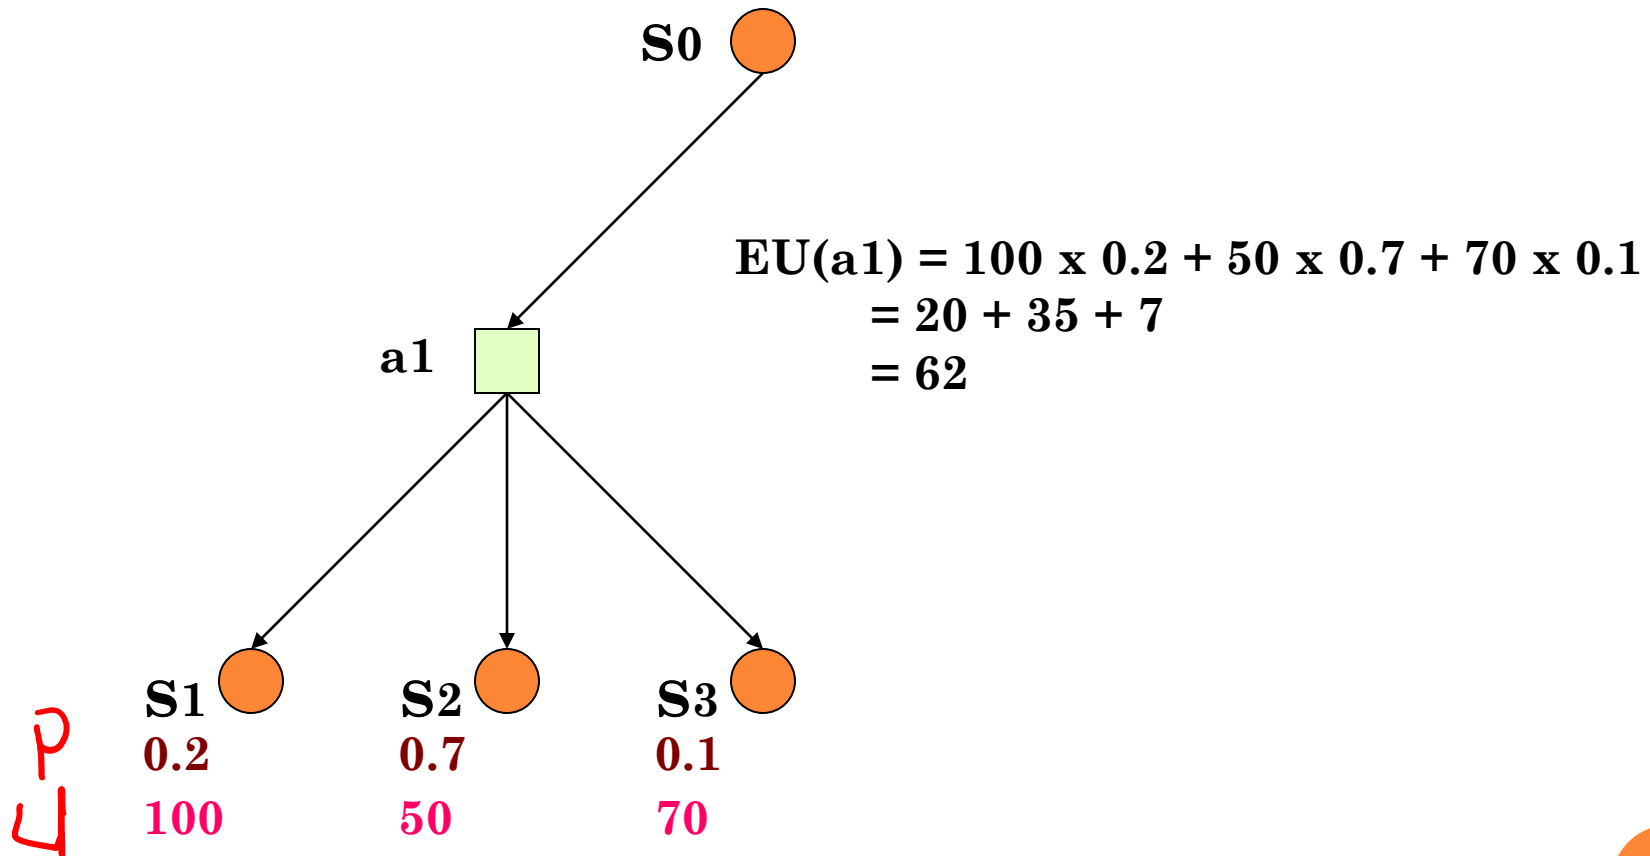
○ Skipped

- Bayesian networks (covered in CS 480 and CS 583)

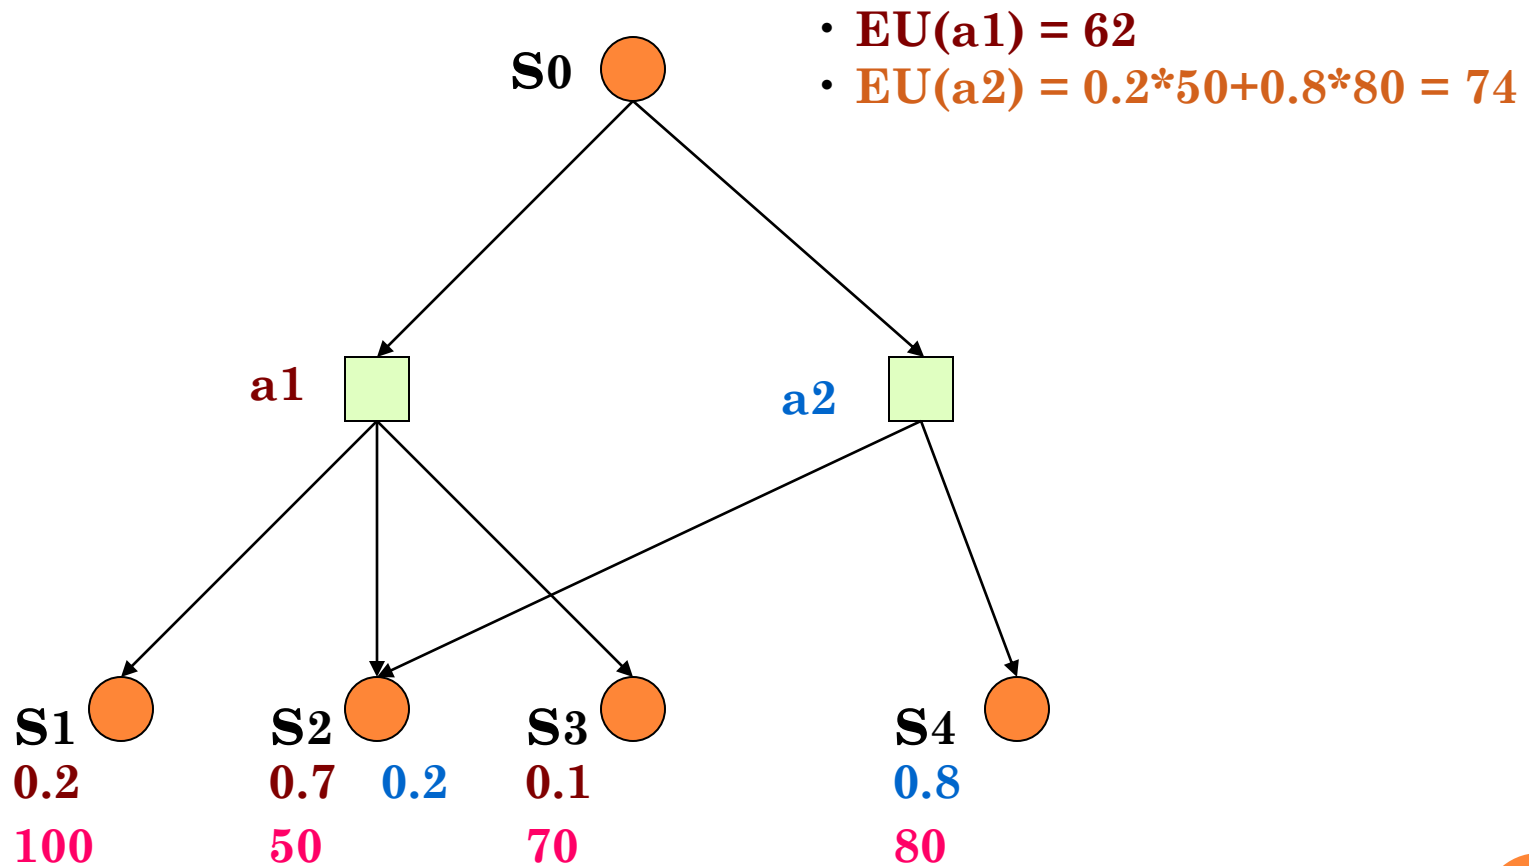
UTILITY THEORY

- Lottery: n possible outcomes with probabilities
 - $[p_1, S_1; p_2, S_2; \dots p_n, S_n]$
 - Each S_i can be an atomic state or another lottery
- **Expected utility of a lottery**
 - $EU([p_1, S_1; p_2, S_2; \dots p_n, S_n]) = \sum_{i=1}^n p_i U(S_i)$

ONE ACTION EXAMPLE



TWO ACTIONS EXAMPLE



UTILITY \neq MONEY

- Most agents prefer more money to less money,
 - But this does not mean money behaves as a utility function
- For example, which lottery would you prefer
 - L_1 : [1, \$1 Million]
 - L_2 : [0.5, \$0; 0.5, \$2.5 Million]
- If money served as a utility function, then you'd prefer L_2 no matter what, but the answer *often* depends on how much money you currently have
 - The utility of money depends on what you prefer
 - If you are short on cash, a little more *certain* money can help
 - If you are already billionaire, you might take the risk
 - Or if you are swimming in debt, you might like to gamble

UTILITY \neq MONEY

- Let's say you currently have \$k and let S_k represent the state of having \$k
- $EU(L_1) = U(S_{k+1M})$
- $EU(L_2) = 0.5 * U(S_k) + 0.5 * U(S_{k+2.5M})$
- The rational choice depends on your preferences for S_k , S_{k+1M} , and $S_{k+2.5M}$
 - i.e, it depends on the values of $U(S_k)$, $U(S_{k+1M})$, and $U(S_{k+2.5M})$
- $U(S_i)$ does not have to be a linear function of i, and for people it often is not

WHICH ACTION TO TAKE?

○ Given

- A probability distribution
- Choice of actions and their effects on the distribution
- Evidence
- A utility function (as a function of states and possibly actions)

○ Find

- Which action maximizes the expected utility?

A SIMPLE SPAM FILTERING EXAMPLE

- Given an email:
 - $P(s | e) = 0.6$
 - $P(\sim s | e) = 0.4$
- Actions:
 - deliverIntoSpamFolder (dS)
 - deliverIntoInboxFolder (dI)
- Utility function:
 - $U(dS, s) = 200$
 - $U(dI, \sim s) = 100$
 - $U(dI, s) = -100$
 - $U(dS, \sim s) = -500$
- Where should this email be delivered and why?

$$t \Rightarrow p \times c + (1-p) \times b$$

UMBRELLA EXAMPLE

$$\neg t \Rightarrow p \times d + (1-p) \times a$$

- Check the weather:

- $P(r) = p$
- $P(\sim r) = 1-p$

- Actions:

- takeUmbrella (t)
- \sim takeUmbrella ($\sim t$)

- Utility function:

- $U(\sim r, \sim t) = a$
- $U(\sim r, t) = b$
- $U(r, t) = c$
- $U(r, \sim t) = d$

- When should you take the umbrella?

VALUE OF INFORMATION

- If I can buy more information to help with my decision, up to how much should I pay for that information?
 - Currency here is in “utility” units
- Value of information
 - Expected utility after the information is acquired
 - Minus
 - Expected utility before the information is acquired
- There is one catch: we do not know the content of the information before we acquire it
 - Solution: take an expectation over the possible outcomes

VALUE OF INFORMATION – ALL BINARY

- Probability of state:
 - $P(s) = p, P(\sim s) = 1-p$
- Actions
 - $a, \sim a$
- Additional info, E
 - $P(s | e) = q, P(\sim s | e) = 1-q$
 - $P(s | \sim e) = r, P(\sim s | \sim e) = 1-r$
 - $P(e) = v, P(\sim e) = 1-v$
- Utilities
 - $U(s, a) = u1$
 - $U(s, \sim a) = u2$
 - $U(\sim s, a) = u3$
 - $U(\sim s, \sim a) = u4$
- What is the value of E ? $VOI(E)$?

NEXT

- Multi-armed bandit
- Markov decision processes
- Reading
 - Chapter 17 of the AI book by Stuart & Russel (<http://aima.cs.berkeley.edu/>)
 - Chapters 2 and 3 of the RL book by Sutton & Barto (<http://www.incompleteideas.net/book/the-book-2nd.html>)

MULTI-ARMED BANDIT

SETUP

- K slot machines
 - Each one is a one-armed bandit
- Unknown reward functions
 - The distribution of rewards for each machine are unknown
- Limited resources
 - Have a limited number of tries (or alternatively, future rewards are discounted)
- Can gather information
 - Each try gives a (potentially) zero/negative reward, but also is useful for information gathering purposes
- Main objective
 - Maximize rewards
- Exploration vs exploitation trade-off
 - How should you balance exploitation (sticking with a machine that looks good) versus exploration (trying new machines)?

NOTATION

- A_t : action at time t
- R_t : reward at time t
- Sequence of actions and rewards:
 $A_1, R_1, A_2, R_2, \dots, A_T, R_T$
- $q_*(a) = E[R_t \mid A_t = a]$
 - Not given; otherwise, the solution is trivial
- $Q_t(a)$: Average reward for action a , up to, excluding, time t
 - Estimated using prior experiences
 - If action a has never been taken before time t , $Q_t(a)$ is initialized to a default value

OPTIMAL STRATEGY

- $\operatorname{argmax}_a q_*(a)$
- At time t , choose the action that has the highest expected reward
- Challenge
 - We do not know $q_*(a)$

GREEDY STRATEGY (EXPLOIT ONLY)

- $\operatorname{argmax}_a Q_t(a)$
- Calculate the average reward for each action, up to time t , and choose the maximum one
- Problems
 - Initial values of $Q_t(a)$ plays a big role
 - It simply tries the best action it has found; purely exploitation
 - Could easily miss other actions that are better but not tried

RANDOM STRATEGY (EXPLORE ONLY)

- Choose an action at random
- Good
 - Explores all the time
- Bad
 - Does not exploit; does not learn

EPSILON GREEDY

- Given an exploration parameter ϵ
- At each step t :
 - With probability ϵ , choose a random action (explore)
 - With probability $1 - \epsilon$, choose the current best action:
 $\operatorname{argmax}_a Q_t(a)$ (exploit)
- $\epsilon = 0$ is fully greedy, and $\epsilon = 1$ is fully random

SIMULATION SETUP




Figure 2.1: An example bandit problem from the 10-armed testbed. The true value $q_*(a)$ of each of the ten actions was selected according to a normal distribution with mean zero and unit variance, and then the actual rewards were selected according to a mean $q_*(a)$, unit-variance normal distribution, as suggested by these gray distributions.

Figure from <http://www.incompleteideas.net/book/the-book-2nd.html>

SIMULATION RESULTS




Figure 2.2: Average performance of ϵ -greedy action-value methods on the 10-armed testbed. These data are averages over 2000 runs with different bandit problems. All methods used sample averages as their action-value estimates.

Figure from <http://www.incompleteideas.net/book/the-book-2nd.html>

UCB

- The upper-confidence bound (UCB) method calculates the upper bound on the mean for each slot and chooses the machine with max value
 - $Q_t(a) + c \sqrt{\frac{\ln(t)}{N_t(a)}}$, where $N_t(a)$ is the number of times the action a is tried and c is the trade-off parameter
- The term in the square root is a measure of uncertainty in $Q_t(a)$; and hence the name upper confidence bound
 - The upper confidence is derived using the Chernoff-Hoeffding bound
- The exploration grows with $\ln(t)$, shrinks with $N_t(a)$
- We've seen this before; where?

SIMULATION RESULTS




Figure 2.4: Average performance of UCB action selection on the 10-armed testbed. As shown, UCB generally performs better than ϵ -greedy action selection, except in the first k steps, when it selects randomly among the as-yet-untried actions.

A RUNNING AVERAGE COMPUTATION

$$\begin{aligned}Q_{n+1} &= \frac{1}{n} \sum_{i=1}^n R_i \\&= \frac{1}{n} \left(R_n + \sum_{i=1}^{n-1} R_i \right) \\&= \frac{1}{n} \left(R_n + (n-1) \frac{1}{n-1} \sum_{i=1}^{n-1} R_i \right) \\&= \frac{1}{n} \left(R_n + (n-1) Q_n \right) \\&= \frac{1}{n} \left(R_n + n Q_n - Q_n \right) \\&= Q_n + \frac{1}{n} \left[R_n - Q_n \right],\end{aligned}$$

Figure from <http://www.incompleteideas.net/book/the-book-2nd.html>

UPDATE RULE

- $Q_{n+1} = Q_n + \frac{1}{n}(R_n - Q_n)$
- $\text{new} = \text{old} + \text{stepSize} * (\text{target} - \text{old})$
- For computing the exact average, stepSize is a function of n

NON-STATIONARY REWARDS

- What if the reward distribution changes over time?
- We'd like to give recent rewards more weight
- A simple approach
 - $Q_{n+1} = Q_n + \alpha(R_n - Q_n)$
- Earlier rewards have lower weights

- $$\begin{aligned} Q_{n+1} &= Q_n + \alpha(R_n - Q_n) \\ &= \alpha R_n + (1 - \alpha)Q_n \\ &= \alpha R_n + (1 - \alpha)[\alpha R_{n-1} + (1 - \alpha)Q_{n-1}] \end{aligned}$$

MARKOV DECISION PROCESSES

PROBLEM SETTING


- The world is represented through states
- At each state, an agent is given 0 (terminal states) or more actions to choose from
- Each action moves the agent, probabilistically, to a state (could be the current state) and results, probabilistically, in a reward (could be zero, negative, positive)
- The agent needs to maximize the sum of the rewards it accumulates over time
- Greedy strategy with respect to immediate rewards often do not work; the agent needs to consider the long-term consequences of its actions

MARKOV DECISION PROCESS

- A sequential decision-making process
- Stochastic environment
- Markov transition model
- Additive rewards
- Applications: planning
 - Search Google
 - <http://www.it.uu.se/edu/course/homepage/aism/st11/MDPApplications1.pdf>

MDP DIAGRAM EXAMPLE

- See OneNote



NOTATION

- $P(s' | s, a)$ Probability of arriving at state s' given we are at state s and take action a
- $R(s, a, s')$ The reward the agent receives when it transitions from state s to state s' via action a
- $\pi(s)$ The action recommended by policy π at state s
- π^* Optimal policy
- $U^\pi(s)$ The expected utility obtained via executing policy π starting at state s
- $U^{\pi^*}(s)$ is often abbreviated as $U(s)$
- $Q^\pi(s, a)$ The expected utility of taking action a at state s
- γ Discount factor $[0, 1]$

RUNNING EXAMPLE




Figure 17.1 (a) A simple, stochastic 4×3 environment that presents the agent with a sequential decision problem. (b) Illustration of the transition model of the environment: the “intended” outcome occurs with probability 0.8, but with probability 0.2 the agent moves at right angles to the intended direction. A collision with a wall results in no movement. Transitions into the two terminal states have reward +1 and -1, respectively, and all other transitions have a reward of -0.04.

Figure from <http://aima.cs.berkeley.edu/figures.pdf>

SOLUTION?

- A fixed action sequence is not the answer due to stochasticity
 - For example, [Up, Up, Right, Right, Right] is not a solution
 - It would be a solution if the environment was deterministic
- A solution must specify the agent should do in any state that the agent might reach
 - This is called a **policy**
- Policy notation: π
 - $\pi(s)$ specifies what action the agent should take at state s
- An **optimal policy** is the one that maximizes the expected utility
 - π^*

RUNNING EXAMPLE

$$\gamma = 1$$

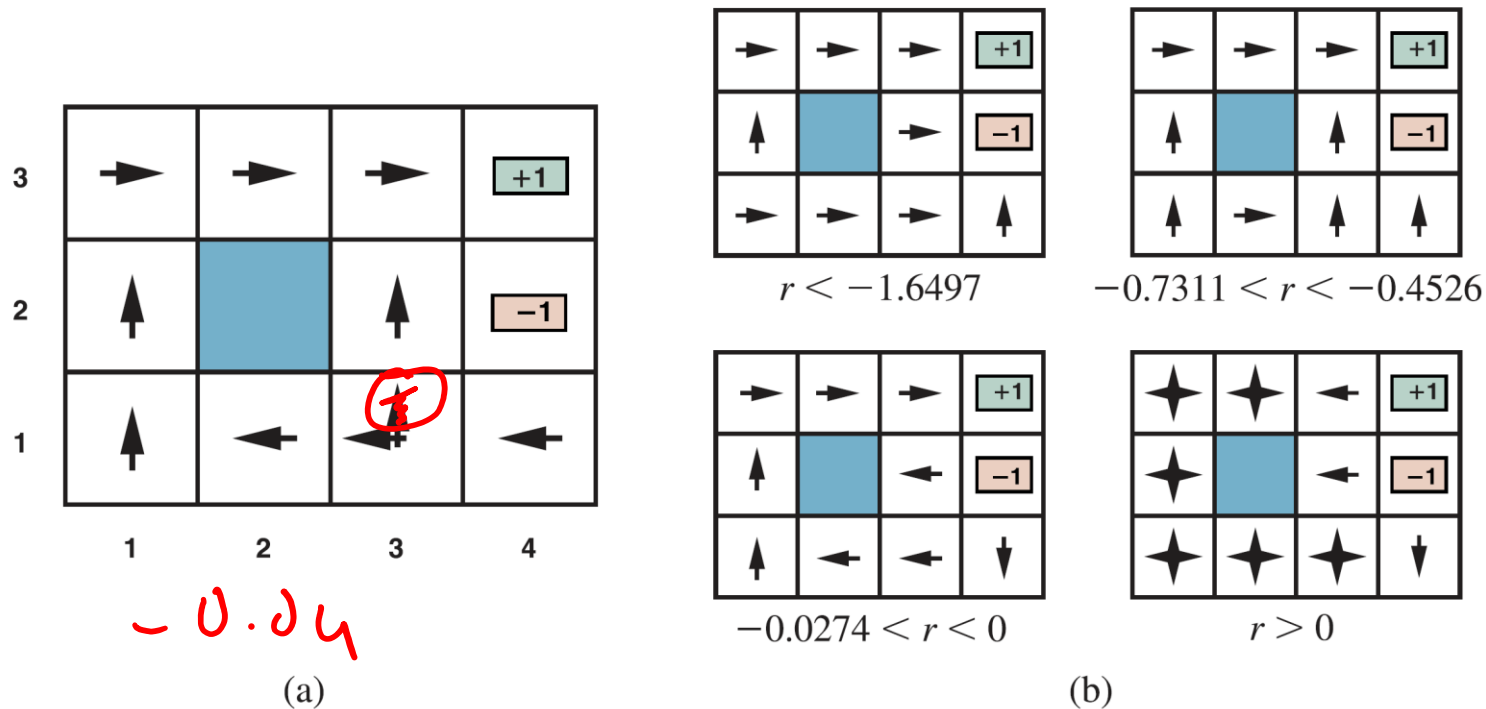


Figure 17.2 (a) The optimal policies for the stochastic environment with $r = -0.04$ for transitions between nonterminal states. There are two policies because in state (3,1) both *Left* and *Up* are optimal. (b) Optimal policies for four different ranges of r .

Figure from <http://aima.cs.berkeley.edu/figures.pdf>

UTILITY OF STATES

- The agent receives a reward at each state
- Utility of a state s given a policy π is the expected reward that the agent will get starting from state s and taking actions according to policy π
- Let S_t denote the state that the agent reaches at time t
- $U^\pi(s) = E[\sum_{t=0}^{\infty} \gamma^t R(S_t, \pi(S_t), S_{t+1})]$
- The expectation is with respect to the transition probabilities

U(s) VS R(s, A, s')

- $R(s, a, s')$ is the short-term immediate reward the agent receives when it transitions from state s to state s' via action a
- $U(s)$ is the long-term cumulative reward from s onward
- $U^\pi(s) = E[\sum_{t=0}^{\infty} \gamma^t R(S_t, \pi(S_t), S_{t+1})]$

BELLMAN EQUATION

$$U^{\pi}(s) = \sum_{s'} P(s'|s, \pi(s)) \times [R(s, \pi(s), s') + \gamma U^{\pi}(s')]$$

BELLMAN OPTIMALITY EQUATION

$$U(s) = \max_{a \in A(s)} \sum_{s'} P(s' | s, a) [R(s, a, s') + \gamma U(s')]$$

THE OPTIMAL POLICY

- The optimal policy is the one that maximizes the expected utility
 - $\pi_s^* = \operatorname{argmax}_{\pi} U^{\pi}(s)$
- Remember that π_s^* is a policy; that is, it recommends an action for each state, regardless of whether it is the starting state or not
- It is optimal when the starting state is s
- When the rewards are discounted, the optimal policy is independent of the start state
 - The optimality of the policy does not depend on the starting state but of course the action sequence depends on the starting state
- True utility of each state is defined as $U^{\pi^*}(s)$ -- the expected rewards the agent will receive if it executes the optimal policy starting at s

ACTION-UTILITY FUNCTION $Q(s, a)$

- $Q^\pi(s, a)$ The expected utility of taking action a at state s and then following policy π
- $Q^\pi(s, a) = \sum_{s'} P(s' | s, a) [R(s, a, s') + \gamma U^\pi(s')]$
- $U^{\pi^*}(s) = \max_{a \in A(s)} Q^{\pi^*}(s, a)$
- $Q^{\pi^*}(s, a) = \sum_{s'} P(s' | s, a) [R(s, a, s') + \gamma \max_{a' \in A(s')} Q^{\pi^*}(s', a')]$
 - Bellman optimality equation for the Q function

$$\pi^*(s) = \operatorname{argmax}_{a \in A(s)} Q^{\pi^*}(s, a)$$

RUNNING EXAMPLE

$$V(T_s) = 0$$

3	0.8516	0.9078	0.9578	A +1
2	0.8016		0.7003	B -1
1	0.7453	0.6953	0.6514	0.4279
	1	2	3	4

Figure 17.3 The utilities of the states in the 4×3 world with $\gamma = 1$ and $r = -0.04$ for transitions to nonterminal states.

Figure from <http://aima.cs.berkeley.edu/figures.pdf>

EXERCISE

- Confirm that the utilities given in the previous slide satisfy the Bellman equations

HOW TO FIND π^*

- $\pi^*(s) = \operatorname{argmax}_{a \in A(s)} \sum_{s'} P(s' | s, a) [R(s, a, s') + \gamma U^{\pi^*}(s')]$
- However, we are not given $U^{\pi^*}(s)$
- Two algorithms for finding optimal policies
 1. Value iteration
 2. Policy iteration

VALUE ITERATION

- $U(s) = \max_{a \in A(s)} \sum_{s'} P(s' | s, a) [R(s, a, s') + \gamma U(s')]$
(Handwritten red annotations: $i+1$ under $U(s)$ and i under $U(s')$)
- n possible states, n Bellman equations, one for each state
- However, these are non-linear equations, due to the max operator
- One approach: iterative
 - Start with an initial guess (could be random)
 - Iterate until convergence

POLICY ITERATION

- Start with an initial policy π_0
- Alternate between
 1. Policy evaluation: given policy π_i , calculate U^{π_i}
 - Can be calculated exactly (linear equations) or iteratively
 2. Policy improvement: Calculate a new MEU policy π_{i+1} , using the utilities calculated in the previous step
- Stop when policy no longer changes

FINDING OPTIMAL POLICIES

- The previous two slides computed U
 - Many people find the concept of U more intuitive than the concept of Q
- Using Q to find the optimal policy makes more sense
 - $\pi^*(s) = \operatorname{argmax}_a Q^{\pi^*}(s, a)$

NEXT

○ Reinforcement learning

- In fact, we already covered many of the fundamentals of RL
 - Value iteration, policy iteration, exploration vs exploitation trade-off
- We are now ready to make the leap from MDPs to RL
- RL can be considered as solving an MDP where the transition and reward dynamics are unknown